

# orderbook

David Kane, Andrew Liu and Khanh Nguyen

## Introduction

The **orderbook** package provides facilities for exploring and visualizing the data associated with an order book. An order book keeps track of the outstanding limit orders for a financial instrument, e.g. a stock. A *limit order* is an order to buy or sell a given quantity of stock at a specified limit price or better. In other words, a limit order will buy (sell) shares at the limit price or lower (higher). Orders will remain in the orderbook until fully executed, so a partial execution where only a fraction of the given quantity is bought or sold will not remove the order. Instead, the order will remain active at the same price, but with the residual quantity of shares.

Consider a simple order book containing five limit orders: sell 150 shares of IBM at \$11.11, sell 150 shares of IBM at \$11.08, buy 100 shares of IBM at \$11.05, buy 200 shares of IBM at \$11.05, and buy 200 shares of IBM at \$11.01.

	Price	Ask Size
	\$11.11	150
	\$11.08	100
300	\$11.05	
200	\$11.01	

Bid Size	Price
----------	-------

Orders on the *bid* (*ask*) side represent orders to buy (sell), and *size* is the number of shares offered at each price level. Here the price levels are \$11.11, \$11.08, \$11.05, and \$11.01. The *inside market* is composed of the *best bid* at \$11.05 (highest bid price) and the *best ask* at \$11.08 (lowest ask price).

(highest bid price) and *best ask* (lowest ask price) at \$11.05 and \$11.08, respectively. The *spread* (\$0.05) is the difference between the best bid and best ask. The *midpoint* (\$11.065) is the average of the best bid and best ask. Every order in the order book has a unique identifier and timestamp indicating the time at which it was accepted into an order book.

There are four types of messages that traders can submit to an exchange: *add*, *cancel*, *cancel/replace*, and *market order*. A trader can *add* a limit order into the order book. She can also *cancel* an order and remove it from the order book.

If a trader wants to update the size of her order, she can issue a *cancel/replace*, which cancels the order, then immediately replaces it with another order at the same price, but with a lower size.

All orders have a *time priority* depending on when they were accepted by the order book. In the case of

two or more orders at the same price level, those orders will be executed in the order they were accepted into the order book. Note that cancel/replace orders can lower the size of an order, but not increase it. Cancel/replace orders maintain the time priority of the original order, so if size increases were allowed traders with orders at the highest time priority for a price level could perpetually increase the size of their order, preventing others from being able to buy or sell stock at that price level.

A trade occurs when a order is executed through a *market order*, which is an order to immediately buy or sell a given quantity of stock at the best available prices. In the above example, suppose that the order to buy 100 shares at \$11.05 was submitted before the order to buy 200 shares at \$11.05. The first order has priority, so if a market order to buy 200 shares is submitted, the first order to buy 100 shares will be completely filled, and the second order to buy 200 shares will only be partially filled.

Now suppose that the market order was to sell 400 shares. Then the first price level would be filled, and 100 shares from the next price level would be filled:

	Price	Ask Size
	\$11.11	150
	\$11.08	100
100	\$11.01	

Bid Size	Price
----------	-------

See Johnson (2010) for further details.

## Examples

NVIDIA is a graphics processing unit and chipset developer with ticker symbol NVDA. Consider the order book for NVDA at a leading electronic exchange on June 8, 2010. We create the orderbook object by specifying the location of our data file.

```
> library(orderbook)

> file <- system.file("data", "sample.txt",
+   package = "orderbook")
> ob <- orderbook(file = file)
> ob <- read.orders(ob, 10000)
> ob
```

An object of class orderbook

```
-----
Current orderbook time:    09:35:02
Message Index:            10,000
Bid Orders:               631
Ask Orders:               1,856
Total Orders:             2,487
```

We read in the first 10,000 messages then show the object. The orderbook time is displayed in 24-hour time, so it is currently 9:35:02 AM. The message index indicates which row in the data file the order book has read through. The display also shows that there are 631 bids and 1,856 asks outstanding, for a total of 2,487 orders. This indicates that many earlier orders have either been cancelled or removed through trades.

```
> summary(ob)
```

Current time is 09:35:02

Ask price levels: 540

Bid price levels: 179

Total price levels: 719

-----

Ask orders: 1,856

Bid orders: 631

Total orders: 2,487

-----

Spread: 0.02

Mid point: 11

-----

Inside market

Best Bid: 11.36

Size: 2,700

Best Ask: 11.38

Size: 400

Using summary the total order information from show is repeated. We see that there are 540 ask and 179 bid price levels, for a total of 719. This indicates that many orders have been submitted at the same price level. The spread is \$0.02, and the midpoint is \$11.37. The inside market is composed of 2,700 shares offered at the best bid of \$11.36 and 400 shares offered at the best ask of \$11.38. This is important because almost all trades occur within the inside market.


```
> display(ob)
```

Current time is 09:35:02

	Price	Ask Size
-----		
	11.42	900
	11.41	1,400
	11.40	1,205
	11.39	1,600
	11.38	400
-----		
2,700	11.36	
1,100	11.35	
1,100	11.34	
1,600	11.33	
700	11.32	
-----		
Bid Size	Price	

display shows the inside market, along with the four next best bid and ask price levels, along with the size at each price level. This gives the user a simple snapshot of the supply and demand in the market.

```
> plot(ob)
```




plot shows a graphical representation of the order book with price levels on the y-axis, and size on the x-axis. The maximum and minimum price levels are 10% above and below the midpoint. Note the large number of shares at \$11.01. It is helpful to know whether the depth at that price level is comprised of a single order, or several. Using the [ method we can view the order information at particular price levels.

```
> ob["11.01"]
```

	price	size	type	time	id
1	11	109	BID	34220988	4403084
2	11	50000	BID	34220988	4403085
3	11	100	BID	34220988	4403086

There is an order for 50,000 shares at the latter price level that accounts for almost all of the size. We can view a plot of the number of orders rather than the number of shares at each price level by specifying type = 'o' when using plot. In the previous plot the maximum and minimum price levels were 10% off from the midpoint, but for this plot we specify a bound of only 3.3% above and below the midpoint.

```
> plot(ob, bounds = 0.033, type = "o")
```



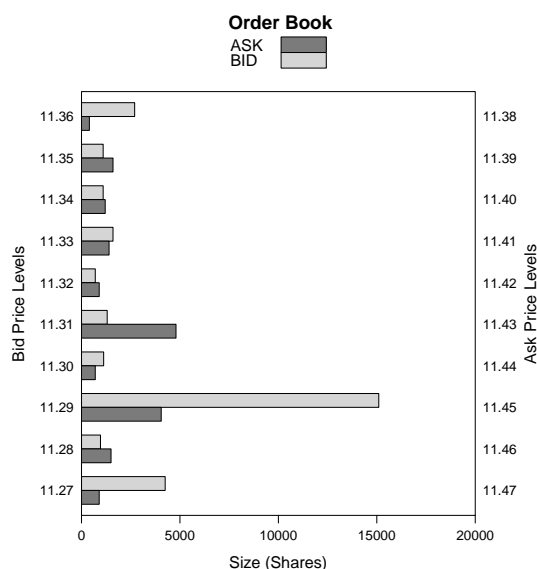
Note the large number of orders at \$11.00. The `[` method returns a `data.frame`, so we can use `nrow` to find out the number of orders at \$11.00.

```
> nrow(ob["11.00"])
```

```
[1] 56
```

There are 56 orders at that price level, which confirms what we see in the plot.

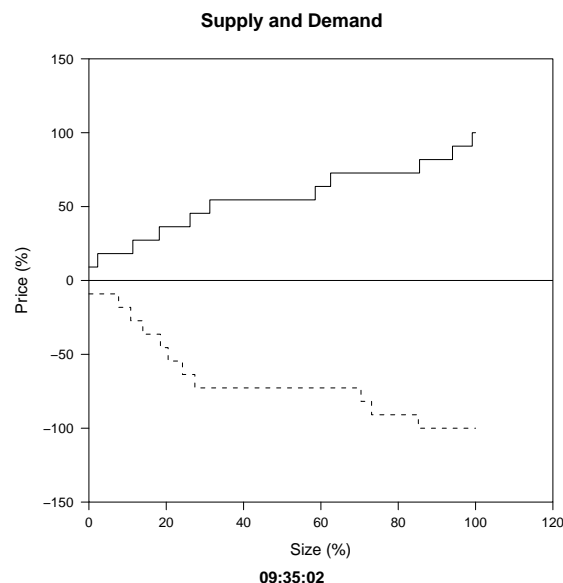
```
> plot(ob, type = "s")
```



Viewing the orderbook with bids on one side and asks on another is useful, but users may want to view them side by side to more directly compare the supply and demand at each price level.

The following plot shows the supply and demand curves for the stock.

```
> plot(ob, bounds = 0.01, type = "sd")
```




This plot shows a downsloping demand curve and an upsloping supply curve. The ask (bid) prices are normalized to be percent above (below) the mid-point. The sizes are normalized by the sum of the sizes across all plotted price levels for each side. See Cao et al. (2009) for more detail.

The user can also view a simple animation of the order book between two times using the `loadanimation` method.

```
> ob <- loadanimation(ob, "9:30:00", "9:31:00")
```

This would generate a Trellis object for each second between 9:30:00 and 9:31:00. The objects would be saved in the R temporary files folder, and their location would be stored in the animation slot within the orderbook object. A single Trellis object is shown below.




The black lines separate individual orders, with the orders nearest the middle y-axis having the highest time priority. The y- and x-axis labels are fixed throughout the animation and the y-axis is chosen so that the user will always see at least five pennies above and below the midpoint price.

`animate` is a simple loop that prints the objects. `pause` specifies the number of seconds to wait in between printing the next Trellis object.

```
> animate(ob, pause = 0.25)
```

The orderbook object also maintains data on trades, which can be visualized using `plot.trade`. This creates a simple bar graph of the number of shares traded at each price level.

```
> plotTrade(ob)
```



Additionally, users can view animations of the order book for a specified number of seconds before and after a trade occurs, or a specified number of messages before and after a trade occurs.

```
> ob <- load.trade.animation(ob, before = 30,
+   after = 30, by = "sec")
```

This would find the next trade from the current order book time, then generate an animation from the 30 seconds before the trade to the 30 seconds after the trade. If the user specified by = 'message' the 30 messages before and after the trade would be shown. In the latter case the current message from the raw data would be printed at the bottom of the plot.

Aside from the ability to retrieve summary statistics and create graphics, **orderbook** can create different orderbook objects for viewing the order book at different times. For example, the user may want to view the order book when the market opens at 9:30:00.

```
> ob <- read.time(ob, "9:30:00")
```

Suppose the user wants to view the last pre-market trade. `previous.trade` finds the first trade that occurred before the current order book time, and then returns an orderbook object at that time. `next.trade` finds the next trade that occurred after the current order book time, and then returns an orderbook object at that time.

```
> ob <- previous.trade(ob)
> ob
```

An object of class orderbook

```
-----
Current orderbook time: 09:34:59
Message Index:         9,958
Bid Orders:            622
Ask Orders:            1,856
Total Orders:          2,478
```

`read.orders` is used to move forwards or backwards in the order book by a specified number of messages. The following will change the state of the orderbook to 50 messages previous to the current message.

```
> ob <- read.orders(ob, n = -50)
> ob
```

An object of class orderbook

```
-----
Current orderbook time: 09:34:55
Message Index:         9,908
Bid Orders:            616
Ask Orders:            1,860
Total Orders:          2,476
```

## Data

Most brokers and exchanges have their own format for transmitting raw order data to customers, so it would be unfeasible for us to write scripts to automatically process that data. Consequently, raw data for an orderbook object must be in the following form:

```
type,time,id,price,size,type
A,31285893,1231884,11.49,200,ASK
R,31295779,1231884,150
T,31295779,1231884,11.49,50
C,31295781,1231884
```

where A, R, T, C mean Add, Replace, Trade, and Cancel, respectively. The first column is the timestamp of the message in milliseconds after midnight of the users timezone, and the second column is the ID of the order. For a cancel/replace the next number is the new size, while for Add and Trade price comes before size, followed by the type of order in the case of Add (BID/ASK).

In this example an order to sell 200 shares at \$11.49 is added to the orderbook, followed by a cancel/replace and a trade several seconds later. Note

that the cancel/replace and the trade have the same timestamp and ID. This is because the orderbook needs to be told the new share size after the trade occurs, as well as information on the trade. It will not adjust the size of a previous order after a trade occurs without an accompanying cancel/replace present. We see that a few milliseconds after the trade the order is entirely cancelled.

## Simulation


**orderbook** supports adding, replacing, and cancelling orders. To add an order, the user needs to specify the price, size, and type. Time and ID are optional, and will default to the maximum time and the maximum ID + 1, respectively. For replacing an order, only ID and size need to be given, and for cancelling an order, only ID is necessary. Market orders are also possible by specifying the size and side (BUY/SELL).

```
> display(ob)
> ob <- add.order(ob, stuff)
> ob <- remove.order(ob, stuff)
> ob <- replace.order(ob, stuff)
> ob <- market.order(ob, 200, "BUY")
> display(ob)
```

Using these tools, the user can write functions to simulate the movement of an order book. In the following example, we consulted Gilles (2006). We simulate 1,000 orders. In each iteration of our simulation there is a 50% chance for a cancel order to be placed, 20% chance for a market order, and 30% chance for a limit order. Orders are cancelled completely randomly, and for a market order there is a 50-50 chance for a buy or sell order to be placed. The size of the market order always corresponds to the size of the best ask or bid at the front of the queue. When a limit order is placed, there is a 50-50 chance for it to be an ask or bid. Then there is a 35% chance for the price to be within the spread, in which case a price is chosen based on a uniform distribution. If the price is determined to be outside of the spread, a price is chosen using a power law distribution. The size follows a log-normal distribution.

```
> ob <- simulate(ob)
```

```
> plot(ob)
```



## Conclusion

The current release of the **orderbook** package is meant to serve as a proof-of-concept. Relatively sophisticated order book analytics are possible using an open source package. The **orderbook** package is part of a collection of packages for performing tests of financial conjectures. See Campbell et al. (2007) and Kane and Enos (2006) for more information on the **backtest** and **portfolio** packages, respectively.

David Kane, Andrew Liu and Khanh Nguyen  
Kane Capital Management  
Cambridge, MA, USA

dave@kanecap.com, Andrew.T.Liu@williams.edu,  
and knguyen@cs.umb.edu

## Bibliography

- K. Campbell, J. Enos, D. Gerlanc, and D. Kane. Backtests. *R News*, 7(1):36–41, April 2007.
- C. Cao, O. Hansch, and X. Wang. The information content of an open limit-order book. *The Journal of Futures Markets*, 29(1):16–41, 2009.
- D. Gilles. *Asynchronous Simulations of a Limit Order Book*. PhD thesis, University of Manchester, 2006.
- B. Johnson. *Algorithmic Trading & DMA: An introduction to direct access trading strategies*. 4Myeloma Press, London, 2010.
- D. Kane and J. Enos. Analysing equity portfolios in r. *R News*, 6(2):13–19, May 2006.